

Julia for Machine Learning – Case Study using miniWeather –

Hyun Kang, Youngsung Kim, and Sarat Sreepathi

Computational Sciences and Engineering Division,
Oak Ridge National Laboratory

2023. 02. 07.

ECPAM



Julia programming for Exascale

Machine learning packages in Julia

- Machine learning packages in Julia
 - A variety of packages are ready in use.
 - Flux
 - FluxMPI
 - Merlin
 - Mocha
 - Knet
 - MLBase
 - ScikitLearn - Julia wrapper for 'ScikitLearn' library
 - TensorFlow - Julia wrapper for 'TensorFlow' library
 - Etc ...
 - Flux.jl
 - A library for machine learning geared towards high-performance production pipelines
 - Many useful tools built in
 - Fully support Julia language
 - NVIDIA GPU support via CUDA.jl

Porting miniWeather-Fortran to Julia

- miniWeather (github.com/mrnorman/miniWeather)
 - Accelerant app mainly developed by Dr. Matthew Norman (ORNL)
 - Simulating weather-like flows for training in parallelizing accelerated HPC architectures
 - Various versions ported to several HPC programming frameworks.
 - MPI (C, Fortran, and C++), OpenACC Offload (C and Fortran), OpenMP Threading (C and Fortran) ...
 - Decent code size (~ 600 SLOC) and well-documented
 - Finite volume spatial discretization (x-z, 2D) and Runge-Kutta time integration
- We used MPI-Fortran version of miniWeather and ported it to Julia.



Porting miniWeather-Fortran to Julia - Porting details

- Porting details
 - Comparison with Fortran codes
 - Array allocation and indices

Fortran

```
allocate(state(1-hs:nx+hs,1-hs:nz+hs, NUM_VARS))
```

Julia (using 'OffsetArray' module)

```
state = OffsetArray(_state,1-HS:NX+HS,1-HS:NZ+HS, NUM_VARS )
```

Column-major, array index starts arbitrary integer

Porting miniWeather-Fortran to Julia - Porting details

- Porting details
 - Comparison with Fortran codes
 - Loops & IF statements

Fortran

```
!Use the fluxes to compute tendencies for each cell
do ll = 1 , NUM_VARS
  do k = 1 , nz
    do i = 1 , nx
      tend(i,k,ll) = -( flux(i,k+1,ll)      &
                        - flux(i,k,ll) ) / dz
      if (ll == ID_WMOM) then
        tend(i,k,ID_WMOM) = tend(i,k,ID_WMOM) &
                              - state(i,k,ID_DENS)*grav
      endif
    enddo
  enddo
enddo
```

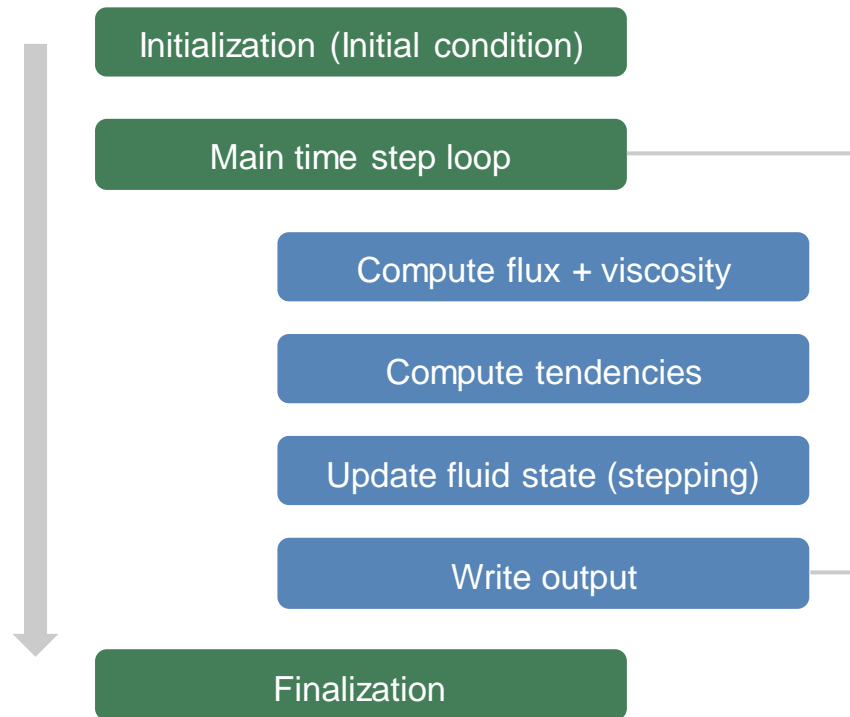
Julia

```
#Use the fluxes to compute tendencies for each cell
for ll in 1:NUM_VARS
  for k in 1:NZ
    for i in 1:NX
      tend[i,k,ll] = -( flux[i,k+1,ll] -
                        flux[i,k,ll] ) / DZ
      if (ll == ID_WMOM)
        tend[i,k,ID_WMOM] = tend[i,k,ID_WMOM] -
                              state[i,k,ID_DENS]*GRAV
      end
    end
  end
end
```

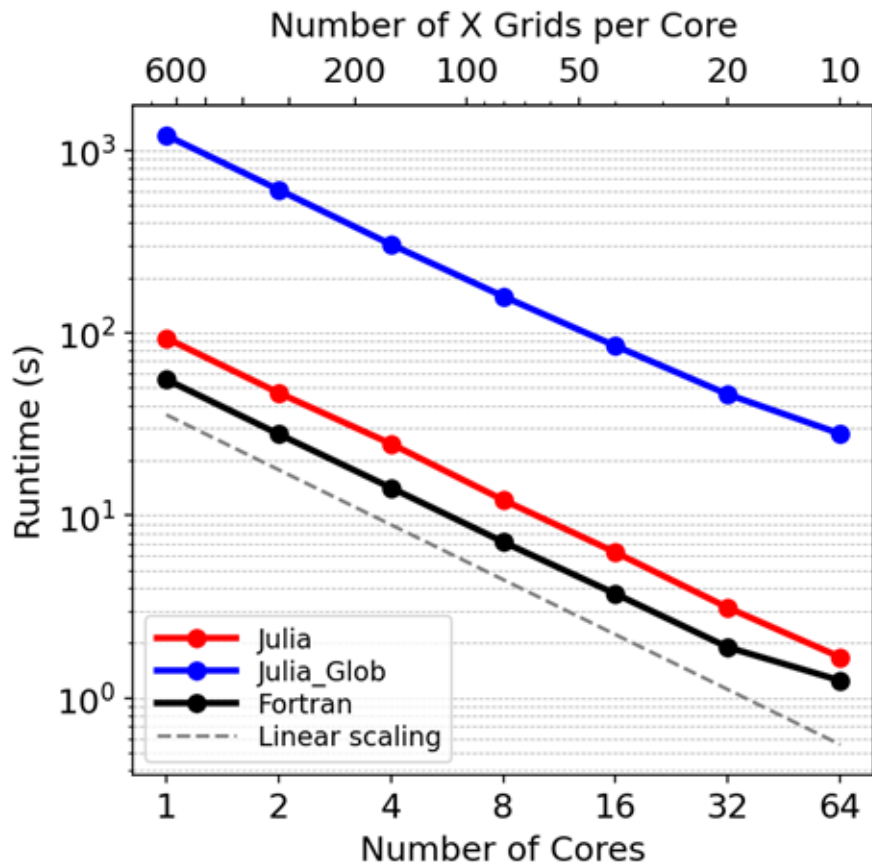
- Julia seems to be Fortran-friendly
 - Ported in 2 days thanks to column-major array and Fortran-style array indexing (OffsetArrays)

Porting miniWeather-Fortran to Julia

- miniWeather model structure



Porting miniWeather-Fortran to Julia - Performance Results



- Fortran is about 1.6x faster than Julia. Wow! Julia is faster than we thought!
- Julia with naively using global variable (Julia_Glob) is about ~13x slower.
 - Be careful when using global variables
- Tested on Cori (Intel KNL) at NERSC
 - On Crusher (a testbed for Frontier), Fortran is about 1.8x faster than Julia.

<https://github.com/grnydawn/jlweather>

Porting miniWeather-Fortran to Julia

- Promising speed of Julia in a simple weather model
⇒ We saw possibilities of Julia for use in HPC area
- Unique benefits as a dynamic language?
 - Fast data processing with convenient use of packages
 - Ex) Huge NetCDF file + custom diagnoses
 - Ex) Packages for data sciences (DataFrames.jl, Pandas.jl (Julia interface for Pandas) ...)
 - Ex) Fast pre+post data processing for ML
 - Seamless integration of machine learning
 - Ex) Replacement of a conventional ML process :
 - Fortran (modeling & data generation & processing) + Python (ML)
→ 2 languages, different coding layer & execution
 - Ex) **Online learning benefitting from fast Julia modeling + convenient ML package use**
 - A common technique used in areas of machine learning where it is computationally infeasible to train over the entire dataset

Machine learning in Julia

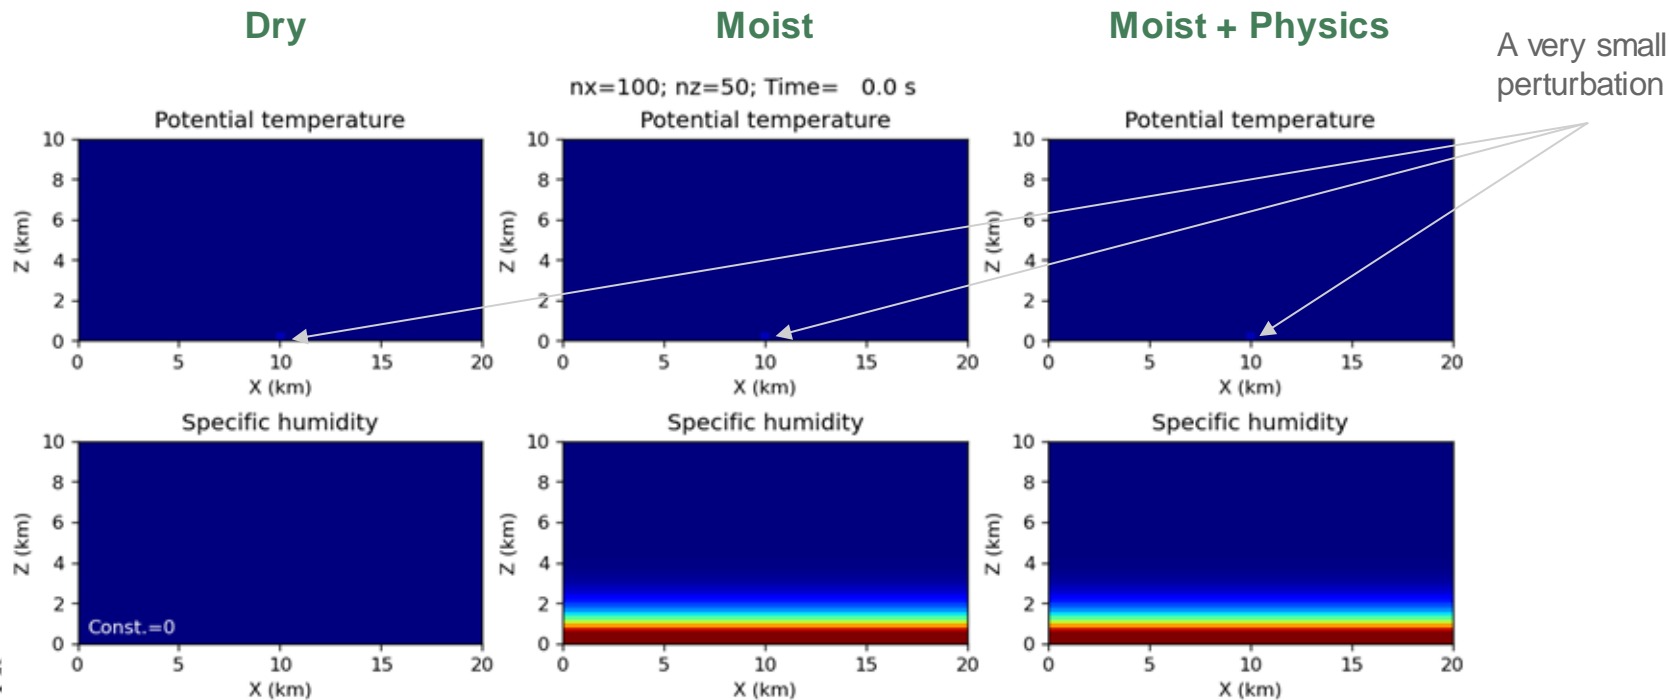
- Machine learning packages in Julia
 - A variety of packages are ready in use.
 - Flux
 - FluxMPI
 - Merlin
 - Mocha
 - Knet
 - MLBase
 - ScikitLearn - Julia wrapper for 'ScikitLearn' library
 - TensorFlow - Julia wrapper for 'TensorFlow' library
 - Etc ...
 - Flux.jl
 - A library for machine learning geared towards high-performance production pipelines
 - Many useful tools built in
 - Fully support Julia language
 - NVIDIA GPU support via CUDA.jl

Machine learning with miniWeather

- Learning simple physics during simulations
 - Physics parameterizations in NWP models usually occupies the second largest portion of a whole runtime.
 - Physics parameterization mimics physical process + sub-grid scale phenomena in numerical weather prediction models
 - E.g., turbulence, precipitation, radiation, etc
 - Emulate an existing physics parameterization to enable faster computation
 - Note: In our demonstration, however, it's a very simple physics schemes. *There is no computational advantage.*
- Coupling parameterized physical processes to miniWeather
 - Simple physics package (Reed and Jablonowski, 2012)
 - Code from DCMIP 2016 (Dynamical Core Model Intercomparison Project (Ullrich et al., 2017))
 - Condensation & Latent heat release (large-scale condensation, no cloud stage)
 - Bulk aerodynamic surface fluxes
 - Boundary-layer mixing

Machine learning using miniWeather

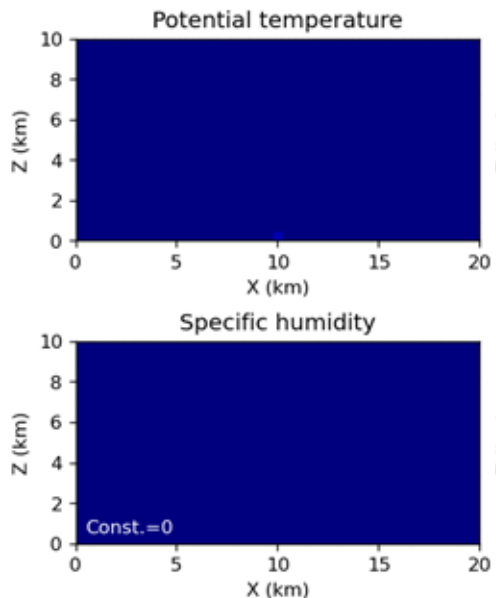
- Simulation results (Dry VS Moist VS Moist + Physics)
 - A very small perturbation in the initial potential temperature
 - Stable atmosphere ($d\theta/dz > 0$)



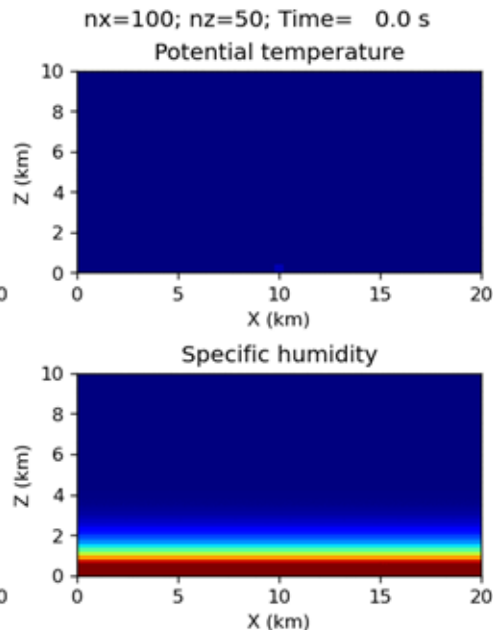
Machine learning using miniWeather

- Simulation results (Dry VS Moist VS Moist + Physics)
 - A very small perturbation in the initial potential temperature
 - Stable atmosphere ($d\theta/dz > 0$)
 - Key process = **Condensation & Latent heat release** → Learn this process


Dry



Moist



Moist + Physics



Machine learning using miniWeather

- Brief introduction of learning process
 - Train data
 - Input (x train data)
 - Normalized T (temperature), Q (specific humidity)
 - Output (y train data)
 - Tendencies (dt/dt, dq/dt) after large-scale condensation process
 - Neural Network using two layers

```
model = Chain(Dense(Qin => Q1 , relu),
                Dense(Q1  => Qout, relu))
```
 - Loss function (MSE)

```
loss(x,y) = Flux.Losses.mse(model(collect(x)), y)
```
 - Optimizer

```
opt = Flux.Descent(lr)
```
 - Training

```
Flux.train!(loss, ps, trainingData, opt)
```

Initialization (Initial condition)

Main time step loop

Compute flux + viscosity

Compute tendencies

Update fluid state (stepping)

Physics parameterization

Training ML model

Finalization

Machine learning using miniWeather

```
function train_model(x,y,etime)

# Parameter sizes
Qin = 2
Q1 = Qin*4
Qout = 2

batchsize = 1
# DataLoader : Performant data loading for deep learning
loader_XY = Flux.Data.DataLoader(
    (x,y),
    batchsize = batchsize,
    shuffle = true)

# Initial time
if etime == 0.0

# Neural network using two layers
model = Chain(Dense(Qin , Q1, relu),
               Dense(Q1 , Qout,relu))

# Activate learning process after 1,500 s model time when flow is active
elseif etime > 1500.0

# Save the model to re-train in next time step
@load "mymodel.bson" model

end
```

```
# Activate learning process after 1,500 s model time when flow is active
if etime == 0.0 || etime > 1500.0

println("Model time = ",etime)

loss(x, y) = Flux.mse(model(x), y) # Our loss function to minimize

lr = 5e-4 # Learning rate
opt = Flux.Descent(lr) # Gradient descent optimizer

epochs = 500 # Number of epochs
trainingLosses = zeros(epochs)

ps = Flux.params(model)

p = Progress(epochs; desc = "Training in progress");

showProgress = true # Display progress bar

# Training loop
@time for ii in 1:epochs
    Flux.train!(loss, ps, loader_XY, opt) # Training the model
    if showProgress
        trainingLosses[ii] = Flux.mean([loss(x,y) for (x,y) in loader_XY])
        next!(p; showvalues = [(:loss, trainingLosses[ii]),
                               (:logloss, log10.(trainingLosses[ii]))], valuecolor = :grey)
    end
end

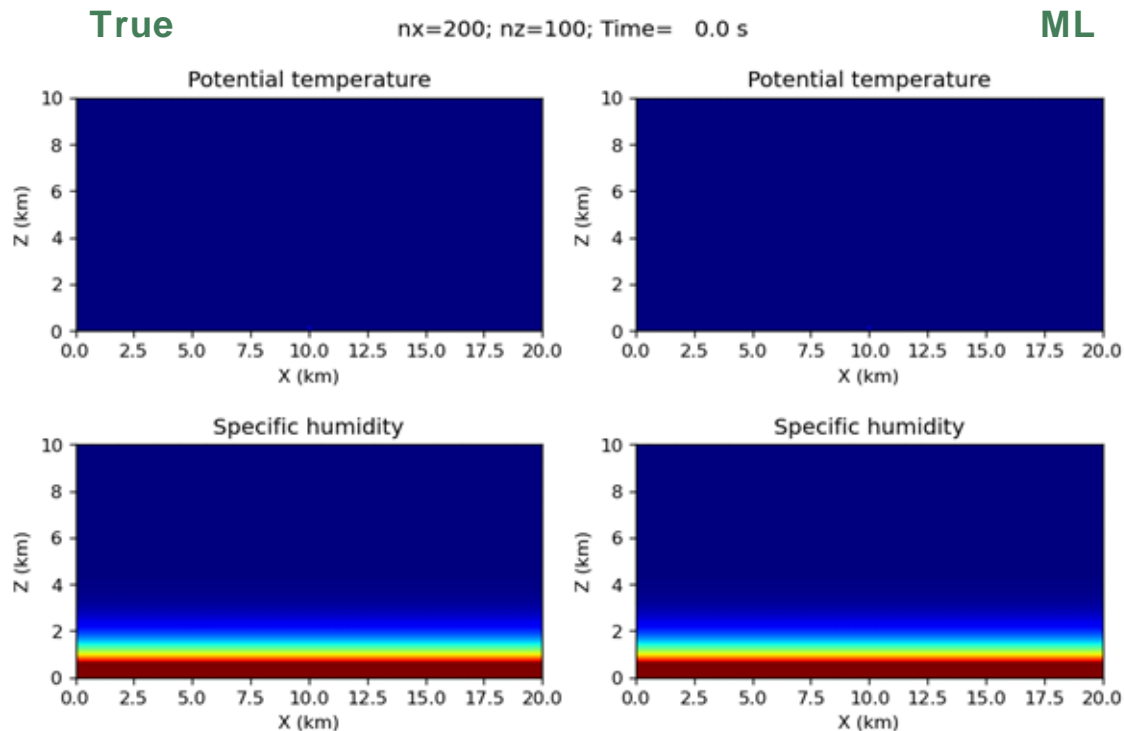
# Save the model to re-train in next time step
@save "mymodel.bson" model

end

end
```

Machine learning using miniWeather

- Results (ML model applied to a 2x higher resolution simulation)
 - ML version mimics the large-scale condensation process.
 - Note: This is a very simple ML model for a demonstration.



Thank you

Machine learning with miniWeather

- Learning simple physics during simulations
 - Physics parameterizations in NWP models usually occupies the second largest portion of a whole runtime.
 - Physics parameterization mimics sub-grid scale phenomena in numerical weather prediction models
 - E.g., cloud, precipitation,
 - Emulate an existing physics parameterization to enable faster computation
 - Note: In our demonstration, however, it's a very simple physics schemes. *There is no computational advantage.*
- Change from dry atmosphere to moist atmosphere in miniWeather
 - Add specific humidity (q) as a tracer
 - Specific humidity affects pressure profile

$$p = C_0 (RT_v)^{\gamma} \quad T_v = (1 + 0.61q) T$$

- Coupling parameterized physical processes to the dynamical core (miniWeather)
 - Simple physics package (Reed and Jablonowski, 2012)
 - Code from DCMIP 2016 (Dynamical Core Model Intercomparison Project (Ullrich et al., 2017))
 - Condensation & Latent heat release (large-scale condensation, no cloud stage)
 - Bulk aerodynamic surface fluxes
 - Boundary-layer mixing

Porting miniWeather-Fortran to Julia - Validation

- Result validations - digit comparison with Fortran results
 - Direct printout from NetCDF output (u-wind comp. at 150 time step)

Fortran

```
5.17481312383945, 5.27095857741744, 4.85877713965263,
```

Julia

```
5.17481312383922, 5.27095857741761, 4.85877713965261,
```

- Globally reduced values (MPI_ALLREDUCE) over 4 cores at 150 time step

Fortran

```
d_te : -5.299941425218241e-04
```

Julia

```
d_te: -5.2999414252196178E-004
```

Around machine precision differences

Porting miniWeather-Fortran to Julia - Porting details

- Porting details
 - Comparison with Fortran codes
 - Subroutine & MPI use

Fortran

```
subroutine reductions( mass , te )
implicit none
real(rp), intent(out) :: mass, te
integer :: i, k, ierr
real(rp) :: r,u,w,th,p,t,ke,ie
real(rp) :: glob(2)
mass = 0
te = 0
do k = 1, nz
  do i = 1, nx
    r = state(i,k,ID_DENS) + hy_dens_cell(k)           ! Density
    u = state(i,k,ID_UMOM) / r                         ! U-wind
    w = state(i,k,ID_WMOM) / r                         ! W-wind
    th = ( state(i,k,ID_RHOT) + hy_dens_theta_cell(k) ) / r ! Potential Temperature
    p = C0*(r*th)**gamma ! Pressure
    t = th / (p0/p)**(rd/cp) ! Temperature
    ke = r*(u*u+w*w) ! Kinetic Energy
    ie = r*cv*t ! Internal Energy
    mass = mass + r *dx*dz ! Accumulate domain mass
    te = te + (ke + r*cv*t)*dx*dz ! Accumulate domain total energy
  enddo
enddo

call mpi_allreduce((/mass,te/),glob,2,mpi_type,MPI_SUM,MPI_COMM_WORLD,ierr)

mass = glob(1)
te = glob(2)
end subroutine reductions
```

Julia

```
function reductions(state::OffsetArray{Float64, 3}, Array{Float64, 3},
  hy_dens_cell::OffsetVector{Float64, Vector{Float64}},
  hy_dens_theta_cell::OffsetVector{Float64, Vector{Float64}})

  local mass, te, r, u, w, th, p, t, ke, le = [zero(Float64) for _ in 1:10]
  glob = Array{Float64}(undef, 2)

  for k in 1:NZ
    for i in 1:NX
      r = state[i,k,ID_DENS] + hy_dens_cell[k]           # Density
      u = state[i,k,ID_UMOM] / r                         # U-wind
      w = state[i,k,ID_WMOM] / r                         # W-wind
      th = ( state[i,k,ID_RHOT] + hy_dens_theta_cell[k] ) / r # Potential Temperature
      p = C0*(r*th)^GAMMA # Pressure
      t = th / (P0/p)^(RD/CP) # Temperature
      ke = r*(u*u+w*w) # Kinetic Energy
      ie = r*CV*t # Internal Energy
      mass = mass + r *DX*DZ # Accumulate domain mass
      te = te + (ke + r*CV*t)*DX*DZ # Accumulate domain total energy
    end
  end

  MPI.Allreduce!(Array{Float64}([mass,te]), glob, +, COMM)

  return glob
end # function
```